

Which side does each output err on, per loss configuration
(shaded half = the output's safe side; same training and test rows)

